

UN-LOCKED ZIG

SYNCHRONIZING CODE WITHOUT LOCKS

Lukas Pietzschmann

October 25th, 2025

Motivation

- Have you ever cooked with others?
 - ▶ It's horrible!
- You need to coordinate who does what and when
 - ▶ Otherwise, you get in each other's way
- Same problem in programming



Generated by Google's Gemini

Concurrent Programming

```
fun Stack(comptime T: type) type {
    return struct {
        backing: std.ArrayList(T),
        pub fn push(self: *Stack, value: T) void {
            // try self.backing.append(value);
            const len = self.backing.items.len;
            self.backing.items.ptr[len] = value;
            self.backing.items.len = len + 1;
        }
    };
}
```

Concurrent Programming

```
fun Stack(comptime T: type) type {
> return struct {
>   > backing: std.ArrayList(T),
>   > lock: std.Thread.Mutex,
>   > pub fn push(self: *Stack, value: T) void {
>     > // try self.backing.append(value);
>     > self.lock.lock();
>     > const len = self.backing.items.len;
>     > self.backing.items.ptr[len] = value;
>     > self.backing.items.len = len + 1;
>     > self.lock.unlock();
>   }
> };
}
```

Beyond Mutual Exclusion

- A Mutex is easy to understand and use
 - ▶ Just grab it and you're safe!
- But for more complex interactions, there are also more complex tools
 - Semaphores
 - Non-blocking Locks
 - Read-Write Locks
 - Reentrant Locks
 - Phases/Barriers
 - ...

The hidden Costs of Locks

Locks seem simple and safe

- But can easily create bottlenecks
- And add additional failure modes
- ▶ Easy to stop thinking about implications

What if we could achieve thread safety without ever forcing a thread to wait?



The hidden Costs of Locks

A selection of additional failure modes

- **Contention**

Multiple threads try to acquire a lock leads to performance degradation.

- **Starvation**

When many threads compete for a lock, some threads may never get it.

- **Priority Inversion**

A lower-priority thread holds a lock needed by a higher-priority thread.

- **Composability**

Locks don't compose well, suggesting the addition of coarser-grained ones.

Agenda

What's
up with
Locks?

Lock-free
Coding

Un-Locking
Zig



Generated by Google's Gemini

How to synchronize without Locks?

The critical section should be so small that no other thread could interrupt it.



Our work horse: Compare and Swap (CAS)

```
fn cas(pointer: *T, expected: T, new: T) bool {  
    if (pointer.* != expected) {  
        return false;  
    }  
    pointer.* = new;  
    return true;  
}
```

“Look at this memory address. If it still contains the value I expected, then – and only then – update it to my new value.”

Building upon CAS

- We can now utilize CAS to actually set a value atomically

```
var current_val: T = atomic_load(prt);
var new_val: T = compute(current_val);
while (!cas(ptr, current_val, new_val)) {
>   current_val = atomic_load(ptr);
>   new_val = compute(current_val);
}
```

- Often, this logic is wrapped into atomic variables
 - ▶ They then provide atomic methods for getting, setting, and updating the value



Levels of Freedom

Obstruction Free

- Weekest Guarantee
- Thread will proceed if all other threads stop

Lock Free

- System-wide progress
- At least one thread makes progress

Wait Free

- Strongest Guarantee
- Per-thread progress

Thread

Data

Suspended
Other Threads

Thread

Data

Other Threads

Thread

Data

Other Threads

Levels of Freedom

Obstruction Free

- Weekest Guarantee
- Thread can't stop if all other threads stop

Lock Free

- System-wide Blocking Algorithms

Wait Free

- Use locks or other blocking primitives
- A thread holding a lock can be suspended, blocking all other threads indefinitely
- ▶ They don't even provide obstruction freedom!

Data

Data

Data

Suspended
Other Threads

Other Threads

Other Threads

Agenda

What's
up with
Locks?

Lock-free
Coding

Un-Locking
Zig



Generated by Google's Gemini

Un-Locking Zig – Atomic Operations

- Our CPUs support atomic operations within their instruction sets
 - test-and-set, fetch-and-increment, compare-and-swap, ...
 - LOCK XCHG, LOCK XADD, LOCK CMPXCHG, ...
- ▶ Lock-free programming is enabled by the hardware itself
- Zig provides access to these atomic operations via built-in functions

```
@cmplxchgWeak(comptime T: type, @cmplxchgStrong(comptime T: type,  
ptr(p*T, *T) != expected) {  
    expected_value: T,  
    new_value: T,  
    success_order: AtomicOrder,  
    fail_order: AtomicOrder  
) ?T  
    ) ?T
```

Un-Locking Zig – Atomic Variables

- **Zig provides atomic types in std.atomic**
 - ▶ std.atomic.Value
- **It provides wrappers around Zig's atomic built-ins**

```
const std = @import("std");
const atomic = std.atomic;

var counter: atomic.Value(u32) = atomic.Value(u32).init(0);
counter.fetchAdd(1, .SeqCst);
```

Un-Locking Zig – Atomic Variables

- **Zig provides atomic types in std.atomic**
 - ▶ std.atomic.Value
- **It provides wrappers around Zig's atomic built-ins**

```
const std = @import("std");
const atomic = std.atomic;

var counter: = atomic.Value(u32) = atomic.Value(u32).init(0);
@atomicRmw(u32, &counter, &Cs.Add, 1, .SeqCst);
```

Let's look at some Code – Push

```
pub fn push(self: *Self, value: T) !void {
> var new_head = try self.allocator.create(Node);
> new_head.* = Node{
> > .value = value,
> > .next = null,
> };
>
> while (true) {
> > const old_head = self.top.load(.acquire);
> > new_head.next = old_head;
> > if (self.top.cmpxchgWeak(old_head, new_head, .release, .←
> acquire) == null) {
> > > return;
> > }
> }
}
```

Let's look at some Code – Push

```
pub fn push(self: *Self, value: T) !void {
> var new_head = try self.allocator.create(Node);
> new_head.* = Node{
> > .value = value,
> > .next = null,
> };
>
> while (true) {
> > const old_head = self.top(.acquire);
> > new_head.next = old_head;
> > if (self.top.cmpxchg(.acquire) == null) {
> > > return;
> > }
> }
> }
```

Why is next not atomic?

/



Let's look at some Code – Pop

```
pub fn pop(self: *Self) ?T {
    while (true) {
        const old_head = self.top.load(.acquire) orelse {
            return null;
        };
        const new_head = old_head.next;
        if (self.top.cmpxchgWeak(old_head, new_head, .release, .←
            acquire) == null) {
            const value = old_head.value;
            self.allocator.destroy(old_head);
            return value;
        }
    }
}
```

Why does CAS take so many parameters?

```
const data = 42;           while (!data_ready) {}  
const data_ready = true;  use(data);
```

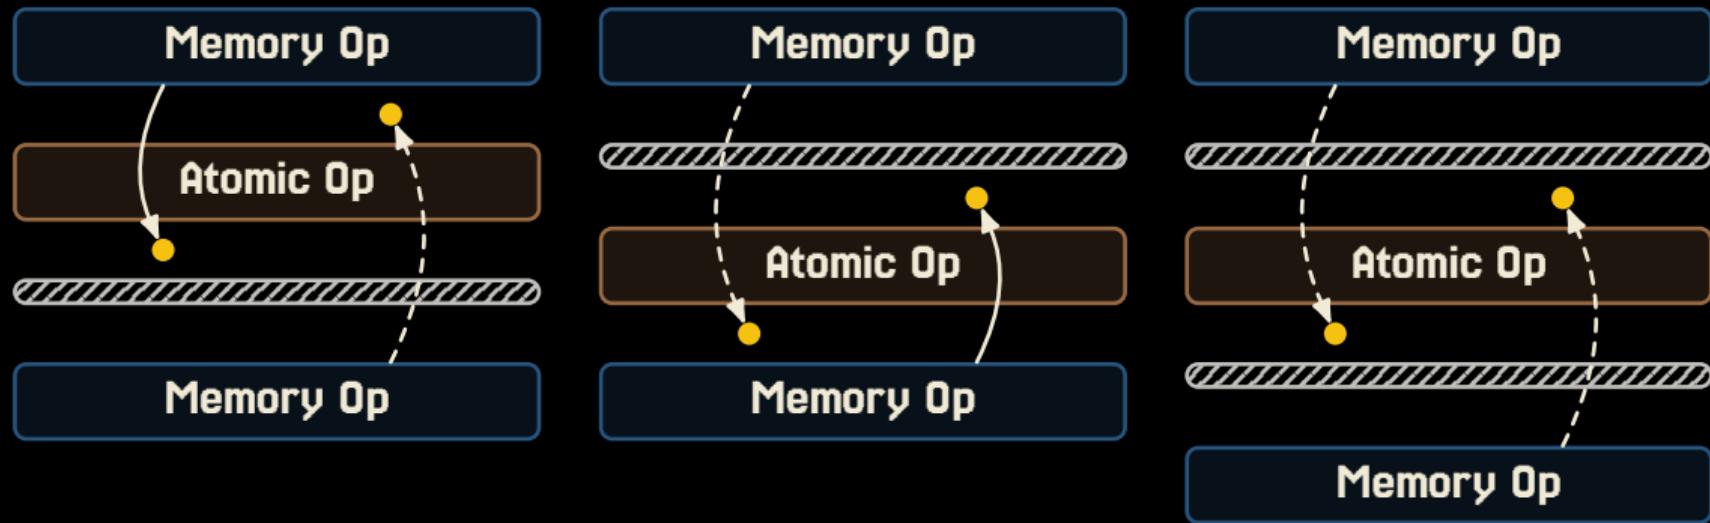
- This seems to work, right?
- ▶ Wrong! There is no clear dependency between data_ready and data
- ▶ The compiler and the CPU might reorder instructions :)
- By introducing memory barriers, we can prevent this reordering
- ▶ Since different operations have different requirements, we need to specify them individually

Bringing Order to Chaos

Acquire

Release

AcqRel



Bringing Order to Chaos

Acquire

Release

AcqRel

```
@cmpxchgWeak(comptime T: type,
    ptr: *T,
    expected_value: T,
    new_value: T,
    success_order: AtomicOrder,
    fail_order: AtomicOrder
) ?T
```

```
@cmpxchgStrong(comptime T: type,
    ptr: *T,
    expected_value: T,
    new_value: T,
    success_order: AtomicOrder,
    fail_order: AtomicOrder
) ?T
```

- The success order is enforced when the the actual and expected values match
- Fail order is enforced when they don't

Memory Op

The Problem with ABA ABA



Generated by Google's Gemini



Generated by Google's Gemini



Generated by Google's Gemini

The Solution to ABA

DCAS

Double CAS; not supported on most hardware.

```
cas(ptr, ep, ap,  
     ver, ev, av);
```

Not to be confused with a wide CAS!

Pointer Tagging

Only delay the problem, but can be practical.

On 8-byte aligned systems, 3 bits are free!
We can just put the version there.

Then we don't need DCAS!

Hazard Pointers

Safe memory reclamation, but complex to implement.

Don't modify the CAS; prevent the "A back to A" part.

Pretty much manual garbage collection.

Agenda

What's
up with
Locks?

Lock-free
Coding

Un-Locking
Zig



Generated by Google's Gemini

When Locks are good enough



Simplicity

- When performance is not critical
- When few threads access a resource a few times

Coordination

- When threads need to wait for each other
- When complex interactions are needed

When Lock-Free shines

Performance

- Better performance under oversubscription
- Better suited for real-time and low-latency systems

Robustness

- No unpredictable blocking delays
- No deadlocks, livelocks, or priority inversions

Conclusion

- Locks have real – often hidden – costs
- Non-blocking algorithms utilize atomic operations to achieve thread safety
- A CAS inside a loop is the building block of many algorithms
- While non-blocking algorithms have actual advantages, they also come with their own challenges

Don't be afraid to use locks, but don't limit yourself to them either!



LUKAS PIETZSCHMANN

Zigtoberfest, October 25th, 2025

lukas@pietzschmann.org